

# Assignment 3

- Dumbvm is a very limited virtual memory system
  - A full TLB leads to a kernel panic
  - Text segment is not read-only
  - Uses fixed segmentation
    - External fragmentation
  - Never reuses physical memory
    - Requires restarting the OS after each test
- Assignment 3 fixes these problems!

# TLB Replacement

- VM related exceptions are handled by *vm\_fault*
- *vm\_fault* performs address translation and loads the virtual address to physical address mapping into the TLB.
  - Iterates through the TLB to find an unused/invalid entry
  - Overwrites the unused entry with the virtual to physical address mapping required by the instruction that generated the TLB exception
- If the TLB is full, call *tlb\_random* to write the entry into a random TLB slot
  - That's it for TLB replacement!
  - Make sure that virtual page fields in the TLB are unique

# Read-Only Text Segment

- Currently, TLB entries are loaded with **TLBLO\_DIRTY** on
  - Pages are therefore read and writeable
- Text segment should be read-only
  - Load TLB entries for the text segment with TLBLO\_DIRTY off
  - `elo &= ~TLBLO_DIRTY;`
- Determine the segment of the fault address by looking at the vbase and vtop addresses

# Read-Only Text Segment

- Unfortunately, this will cause `load_elf` to throw a `VM_FAULT_READONLY` exception
  - It is trying to write to a memory location that is read-only
- We must instead load TLB entries with `TLBLO_DIRTY` on until `load_elf` has completed.
  - Consider adding a flag to **`struct addrspace`** to indicate whether or not `load_elf` has completed
  - When `load_elf` completes, flush the TLB, and ensure that all future TLB entries for the text segment has `TLBLO_DIRTY` off

# Read-Only Text Segment

- Writing to read-only memory address will lead to a `VM_FAULT_READONLY` exception
  - This will currently cause a kernel panic
- Instead of panicing, your VM system should kill the process
  - Modify `kill_curthread` to kill the current process
    - Very similar to `sys__exit`. However, the exit code/status will be different
    - Consider which signal number this will trigger (look at the beginning of `kill_curthread`)

# Managing Memory

Physical memory

0x0

A large, empty rectangle with a thin blue border, representing a block of physical memory. It is positioned centrally below the title and above the address label.

# Managing Memory

Physical memory

During bootstrap, the kernel allocates memory by calling `getppages`, which in turn calls `ram_stealmem(pages)`.

`ram_stealmem` just allocates pages without providing any mechanism to release pages (see `free_kpages`)

We want to manage our own memory after bootstrap

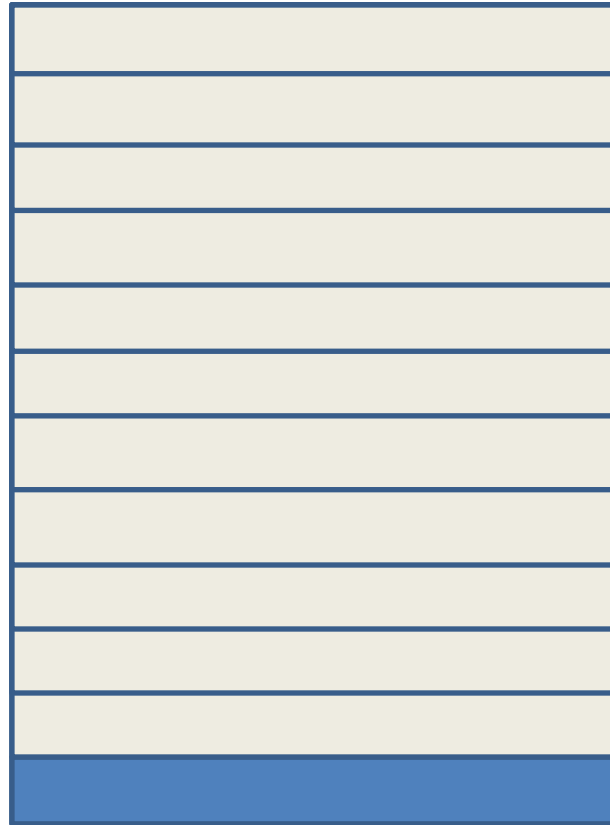
0x0



Memory for bootstrap

# Managing Memory

Physical memory



In `vm_bootstrap`, call `ram_getsize` to get the remaining physical memory in the system. Do not call `ram_stealmem` again!

Logically partition the memory into fixed size frames. Each frame is `PAGE_SIZE` bytes.

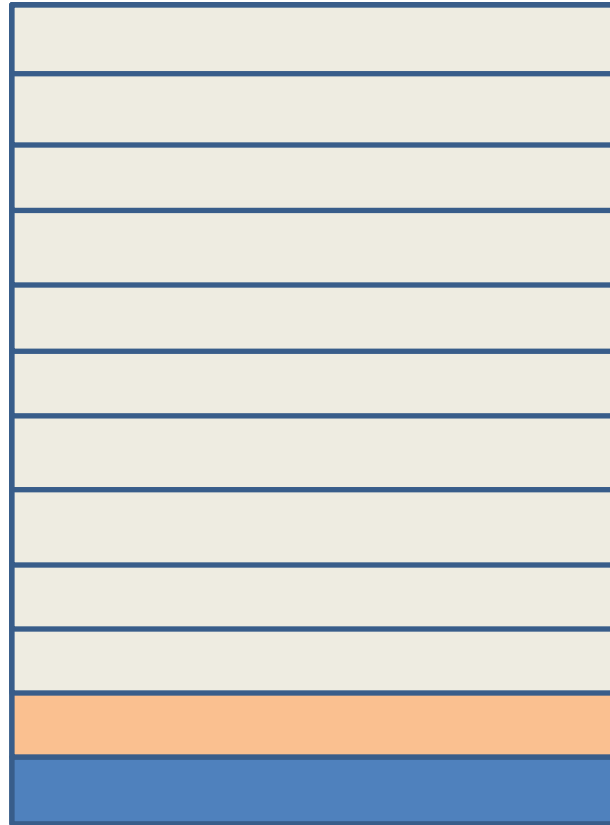
Keep track of the status of each frame (core-map data structure).

Memory for bootstrap



# Managing Memory

Physical memory



Core-map  
Memory for bootstrap

Where should we store the core-map data structure?

Store it in the start of the memory returned by `ram_getsize`. The frames should start after the core-map data structure.

The core-map should track which frames are used and available. It should also keep track of contiguous memory allocations (why?)

# Alloc and Free

- `alloc_kpages(int npages):`
  - Allocate frames for both `kmalloc` and address spaces
  - Frames need to be contiguous
- `free_kpages(vaddr_t addr):`
  - Free pages allocated with `alloc_kpages`
  - We don't specify how many pages we need to free. It should free the same number of pages that was allocated.
  - Update core-map to make frames available after `free_kpages`
- Consider adding some information in the core-map to help determine the number of pages that need to be freed
  - E.g. If 4 contiguous frames were allocated using `alloc_kpages`, then store 4 in the core-map entry for the start of the four frames

# Page Tables

- In order to avoid external fragmentation, we need to introduce paging
- New VM system combines segmentation with paging
- Three segments:
  - Text (read-only)
  - Data
  - Stack
- Create a page table for each segment
  - Each page table entry should include the frame number

# Page Tables

- In dumbvm, struct addrspace has the following fields:
  - vaddr\_t as\_vbase1
  - paddr\_t as\_pbase1
  - size\_t as\_npages1
  - vaddr\_t as\_vbase2
  - paddr\_t as\_pbase2
  - size\_t as\_npages2
  - paddr\_t as\_stackpbase
- With segmentation and paging, what fields should we have instead?
  - Replace pbase with page table

# Page Tables

- `as_create`:
  - Initialize address space data structures
- `as_define_region`:
  - A region is essentially a segment
  - Allocate (`kmalloc`) and initialize the page table for the specified segment
  - Size of the segment is a parameter of `as_define_region`
    - Because we perform preloading, segment size will never grow!
    - Size of the page table is based on the segment size
  - Setup the read/write permissions for this segment
  - Optionally have permissions per page

# Page Tables

- `as_prepare_load`:
  - Pre-allocate frames for each page in the segment
  - Frames do not need to be contiguous
  - Allocate each frame one at a time
- `as_define_stack`:
  - Always allocate `NUM_STACK_PAGES` for the stack
  - Need to create a page table for the stack
  - Need to allocate frames for the stack
    - `as_prepare_load` only allocates frames for segments that were defined by `load_elf`
    - Stack segment is not defined by `load_elf`

# Page Tables

- `as_copy`:
  - Call *as\_create* to create the address space
  - Create segments based on old address space
  - Allocate frames for the segments
  - `memcpy` frames from the old address space to the frames in the new address space
- `as_destroy`:
  - Call `free_kpages` on the frames for each segment
  - `kfree` the page tables

# User Address/Kernel Virtual Address/ Physical Address

- Remember that you are always working with virtual addresses
  - Only use physical addresses when loading entries in the TLB
  - Virtual addresses are converted either by the TLB or by the MMU directly
- Addresses below 0x80000000 are user space addresses that are TLB mapped
- Addresses between 0x80000000 and 0xa0000000 are kernel virtual addresses that are converted by the MMU directly
  - Kernel virtual address – 0x80000000 = physical address
- kmalloc always returns a kernel virtual address. Do not use kmalloc to allocate frames