

Semaphore P() implementation

```
1  P(struct semaphore* sem) {
2      spinlock_acquire(&sem->sem_lock);
3      while (sem->sem_count == 0) {
4          wchan_lock(sem->sem_wchan);
5          spinlock_release(&sem->sem_lock);
6          wchan_sleep(sem->sem_wchan);
7          spinlock_acquire(&sem->sem_lock);
8      }
9      sem->sem_count--;
10     spinlock_release(&sem->sem_lock);
11 }
```

What if we swapped lines 4 and 5?

```
1  P(struct semaphore* sem) {
2      spinlock_acquire(&sem->sem_lock);
3      while (sem->sem_count == 0) {
4          wchan_lock(sem->sem_wchan);
5          spinlock_release(&sem->sem_lock);
6          wchan_sleep(sem->sem_wchan);
7          spinlock_acquire(&sem->sem_lock);
8      }
9      sem->sem_count--;
10     spinlock_release(&sem->sem_lock);
11 }
```

Incorrect implementation

Semaphore count is initially 0

Thread 1

Thread 2

Calls P()



spinlock_release(...)

```
1 P(struct semaphore* sem) {
2     spinlock_acquire(&sem->sem_lock);
3     while (sem->sem_count == 0) {
4         spinlock_release(&sem->sem_lock);
5         wchan_lock(sem->sem_wchan);
6         wchan_sleep(sem->sem_wchan);
7         spinlock_acquire(&sem->sem_lock);
8     }
9     sem->sem_count--;
10    spinlock_release(&sem->sem_lock);
11 }
```

```
1 V(struct semaphore* sem) {
2     spinlock_acquire(&sem->sem_lock);
3     sem->sem_count++;
4     wchan_wakeone(sem->sem_wchan);
5     spinlock_release(&sem->sem_lock);
6 }
```

Incorrect implementation

Semaphore count is now 1

Thread 1

Thread 2

Calls P()

spinlock_release(...)

Either context switch or
just very slow execution
(e.g., pipeline stall,
cache miss, etc.)

Calls V()

Increments count to 1
Calls wchan_wakeone

```
1 P(struct semaphore* sem) {
2     spinlock_acquire(&sem->sem_lock);
3     while (sem->sem_count == 0) {
4         spinlock_release(&sem->sem_lock);
5         wchan_lock(sem->sem_wchan);
6         wchan_sleep(sem->sem_wchan);
7         spinlock_acquire(&sem->sem_lock);
8     }
9     sem->sem_count--;
10    spinlock_release(&sem->sem_lock);
11 }
```

```
1 V(struct semaphore* sem) {
2     spinlock_acquire(&sem->sem_lock);
3     sem->sem_count++;
4     wchan_wakeone(sem->sem_wchan);
5     spinlock_release(&sem->sem_lock);
6 }
```

Incorrect implementation

Semaphore count is now 1

```
1 P(struct semaphore* sem) {
2     spinlock_acquire(&sem->sem_lock);
3     while (sem->sem_count == 0) {
4         spinlock_release(&sem->sem_lock);
5         wchan_lock(sem->sem_wchan);
6         wchan_sleep(sem->sem_wchan);
7         spinlock_acquire(&sem->sem_lock);
8     }
9     sem->sem_count--;
10    spinlock_release(&sem->sem_lock);
11 }
```

```
1 V(struct semaphore* sem) {
2     spinlock_acquire(&sem->sem_lock);
3     sem->sem_count++;
4     wchan_wakeone(sem->sem_wchan);
5     spinlock_release(&sem->sem_lock);
6 }
```

Thread 1

Calls P()



spinlock_release(...)



wchan_lock(...)
wchan_sleep(...)

Either context switch or
just very slow execution
(e.g., pipeline stall,
cache miss, etc.)

Thread 2

Calls V()



Increments count to 1
Calls wchan_wakeone

Incorrect implementation

Semaphore count is now 1

Thread 1

Thread 2

Calls P()

`spinlock_release(...)`

Either context switch or
just very slow execution
(e.g., pipeline stall,
cache miss, etc.)

`wchan_lock(...)`
`wchan_sleep(...)`

Calls V()

Increments count to 1
Calls `wchan_wakeone`

**Thread 1 misses the wake up signal and
blocks on P() even though count is 1.**

```
1 P(struct semaphore* sem) {
2     spinlock_acquire(&sem->sem_lock);
3     while (sem->sem_count == 0) {
4         spinlock_release(&sem->sem_lock);
5         wchan_lock(sem->sem_wchan);
6         wchan_sleep(sem->sem_wchan);
7         spinlock_acquire(&sem->sem_lock);
8     }
9     sem->sem_count--;
10    spinlock_release(&sem->sem_lock);
11 }
```

```
1 V(struct semaphore* sem) {
2     spinlock_acquire(&sem->sem_lock);
3     sem->sem_count++;
4     wchan_wakeone(sem->sem_wchan);
5     spinlock_release(&sem->sem_lock);
6 }
```

Correct implementation

Semaphore count is initially 0

Thread 1

Thread 2

Calls P()



wchan_lock(sem->sem_wchan)
spinlock_release(...)

```
1 P(struct semaphore* sem) {  
2     spinlock_acquire(&sem->sem_lock);  
3     while (sem->sem_count == 0) {  
4         wchan_lock(sem->sem_wchan);  
5         spinlock_release(&sem->sem_lock);  
6         wchan_sleep(sem->sem_wchan);  
7         spinlock_acquire(&sem->sem_lock);  
8     }  
9     sem->sem_count--;  
10    spinlock_release(&sem->sem_lock);  
11 }
```

```
1 V(struct semaphore* sem) {  
2     spinlock_acquire(&sem->sem_lock);  
3     sem->sem_count++;  
4     wchan_wakeone(sem->sem_wchan);  
5     spinlock_release(&sem->sem_lock);  
6 }
```

Correct implementation

Semaphore count is now 1

Thread 1

Thread 2

Calls P()

wchan_lock(sem->sem_wchan)
spinlock_release(...)

Either context switch or
just very slow execution
(e.g., pipeline stall,
cache miss, etc.)

Calls V()

Increments count to 1.
Calls *wchan_wakeone*
Spins inside
wchan_wakeone while
trying to acquire the
channel lock.

```
1 P(struct semaphore* sem) {
2     spinlock_acquire(&sem->sem_lock);
3     while (sem->sem_count == 0) {
4         wchan_lock(sem->sem_wchan);
5         spinlock_release(&sem->sem_lock);
6         wchan_sleep(sem->sem_wchan);
7         spinlock_acquire(&sem->sem_lock);
8     }
9     sem->sem_count--;
10    spinlock_release(&sem->sem_lock);
11 }
```

```
1 V(struct semaphore* sem) {
2     spinlock_acquire(&sem->sem_lock);
3     sem->sem_count++;
4     wchan_wakeone(sem->sem_wchan);
5     spinlock_release(&sem->sem_lock);
6 }
```


Correct implementation

Semaphore count is now 1

Thread 1

Thread 2

Calls P()

Calls V()

`wchan_lock(sem->sem_wchan)`
`spinlock_release(...)`

Either context switch or
just very slow execution
(e.g., pipeline stall,
cache miss, etc.)

Increments count to 1.
Calls `wchan_wakeone`
Spins inside
`wchan_wakeone` while
trying to acquire the
channel lock.

`wchan_sleep(...)`

`wchan_wakeone(...)`
`spinlock_release(...)`

`spinlock_acquire(...)`

`wchan_sleep` releases the
channel lock, which allows
`wchan_wakeone` to complete.
This in turn wakes up thread 1.

```
1 P(struct semaphore* sem) {
2     spinlock_acquire(&sem->sem_lock);
3     while (sem->sem_count == 0) {
4         wchan_lock(sem->sem_wchan);
5         spinlock_release(&sem->sem_lock);
6         wchan_sleep(sem->sem_wchan);
7         spinlock_acquire(&sem->sem_lock);
8     }
9     sem->sem_count--;
10    spinlock_release(&sem->sem_lock);
11 }
```

```
1 V(struct semaphore* sem) {
2     spinlock_acquire(&sem->sem_lock);
3     sem->sem_count++;
4     wchan_wakeone(sem->sem_wchan);
5     spinlock_release(&sem->sem_lock);
6 }
```

However, our implementation does not provide FIFO ordering.

Thread 1

Thread 2

Thread 3

```
1 P(struct semaphore* sem) {
2     spinlock_acquire(&sem->sem_lock);
3     while (sem->sem_count == 0) {
4         wchan_lock(sem->sem_wchan);
5         spinlock_release(&sem->sem_lock);
6         wchan_sleep(sem->sem_wchan);
7         spinlock_acquire(&sem->sem_lock);
8     }
9     sem->sem_count--;
10    spinlock_release(&sem->sem_lock);
11 }
```

Calls P()



wchan_sleep(...)

In the ready queue

```
1 V(struct semaphore* sem) {
2     spinlock_acquire(&sem->sem_lock);
3     sem->sem_count++;
4     wchan_wakeone(sem->sem_wchan);
5     spinlock_release(&sem->sem_lock);
6 }
```

However, our implementation does not provide FIFO ordering.

```
1 P(struct semaphore* sem) {
2     spinlock_acquire(&sem->sem_lock);
3     while (sem->sem_count == 0) {
4         wchan_lock(sem->sem_wchan);
5         spinlock_release(&sem->sem_lock);
6         wchan_sleep(sem->sem_wchan);
7         spinlock_acquire(&sem->sem_lock);
8     }
9     sem->sem_count--;
10    spinlock_release(&sem->sem_lock);
11 }
```

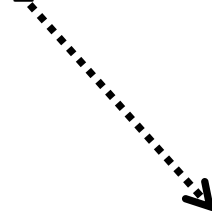
```
1 V(struct semaphore* sem) {
2     spinlock_acquire(&sem->sem_lock);
3     sem->sem_count++;
4     wchan_wakeone(sem->sem_wchan);
5     spinlock_release(&sem->sem_lock);
6 }
```

Thread 1

Calls P()



wchan_sleep(...)



Thread 2

Calls V()



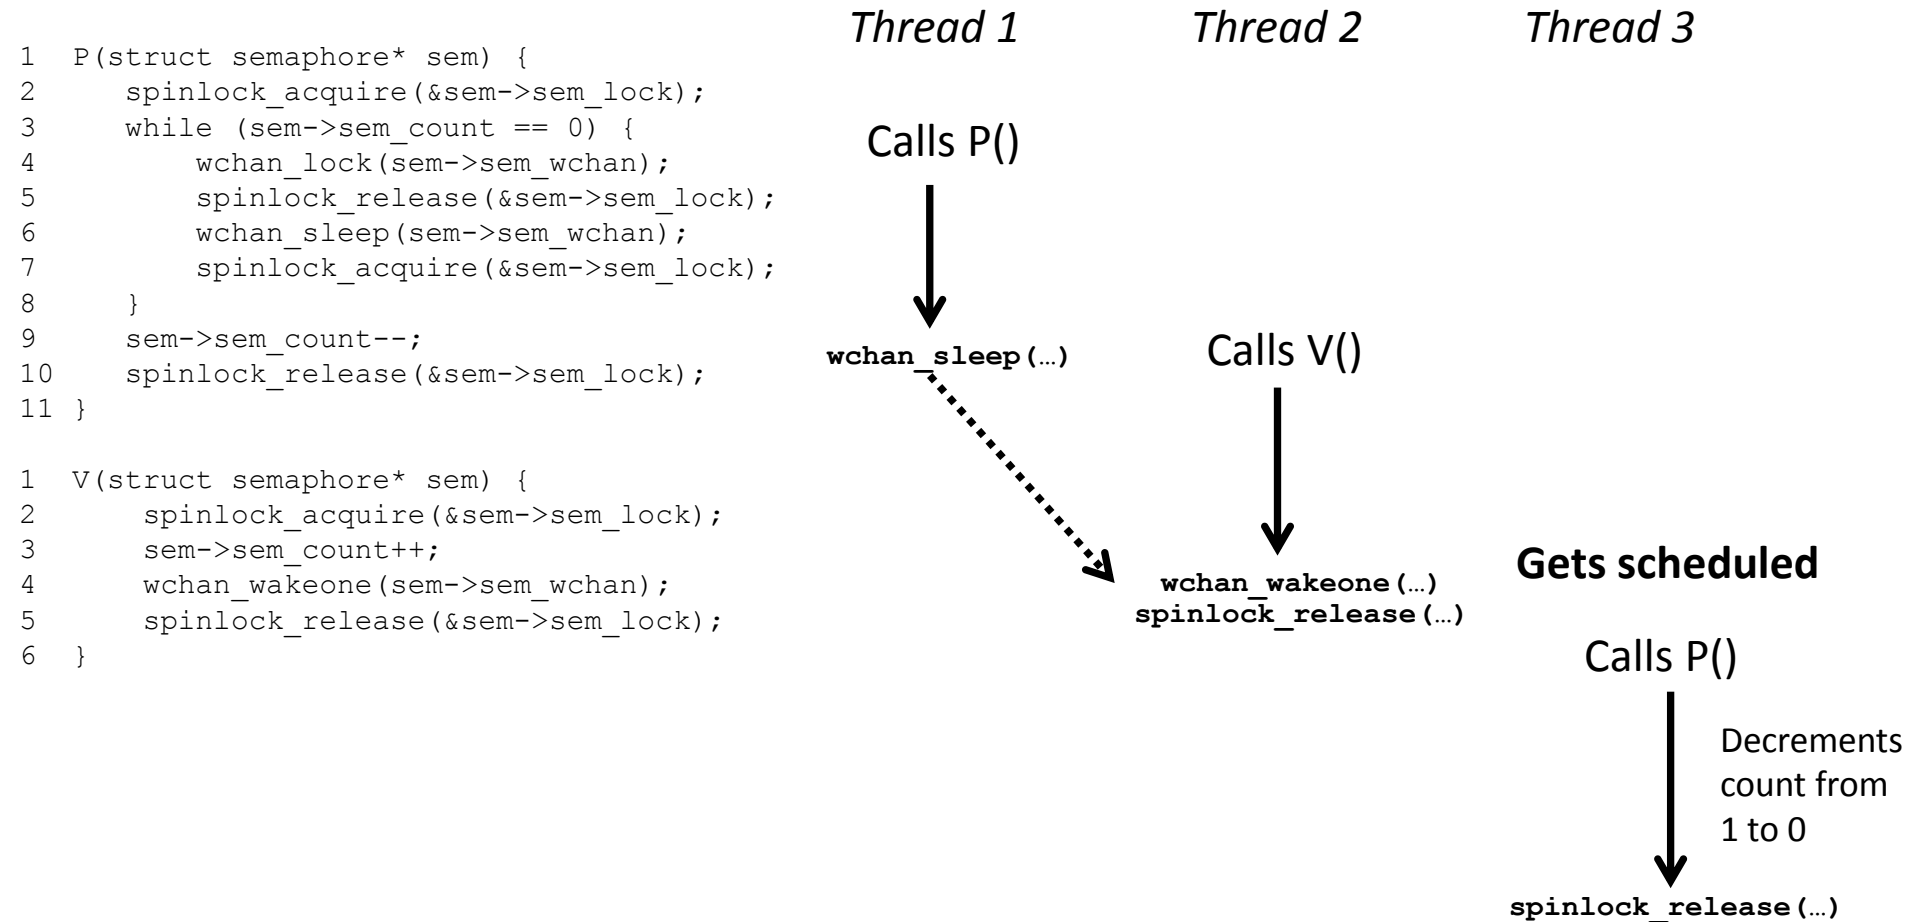
wchan_wakeone(...)
spinlock_release(...)

Thread 3

In the ready queue

Wakes thread 1 up.
Thread 1 added to
the end of the
ready queue

However, our implementation does not provide FIFO ordering.



However, our implementation does not provide FIFO ordering.

