

- **Mutual Exclusion vs. Synchronization**

- Mutual exclusion: Ensures that a shared object can only be accessed by one thread at a time
 - Provided by: Spinlock, binary semaphore, blocking lock
- Synchronization: Establishes timing/order relationships among threads. For example:
 - Thread A should always run after thread B
 - Stage $i + 1$ can't start until stage i has completed
 - Only K threads should be running at the same time
 - Producer/consumer problem with a bounded buffer
 - A thread should only execute if a particular condition is true
 - Wait until the stock price is above \$X before selling the stock
 - Determining when a car/thread should be entering an intersection
 - Provided by: Counting semaphore and condition variable

- **Wait Channel Review**

- Wait channel is a structure defined as follows:

```
struct wchan {  
    const char *wc_name; /* name for this channel */  
    struct threadlist wc_threads; /* list of waiting threads */  
    struct spinlock wc_lock; /* lock for mutual exclusion */  
};
```

- *wchan_sleep* adds the current thread to the thread list
- *wchan_wakeone* moves the first thread from the list to the end of the ready queue
- *wchan_wakeall* moves all of the threads from the list to the ready queue
- They must acquire and release *wc_lock*

- *wchan_wakeone* and *wchan_wakeall* acquire and release *wc_lock* on their own:

```
void
wchan_wakeone(struct wchan *wc)
{
    struct thread *target;

    /* Lock the channel and grab a thread from it */
    spinlock_acquire(&wc->wc_lock);
    target = threadlist_remhead(&wc->wc_threads);
```

- *wchan_sleep* does not acquire *wc_lock* on its own
 - Must explicit call *wchan_lock* before calling *wchan_sleep*

- This is because:
 - Wait channels are almost never used on their own
 - Too easy to miss a wakeup signal due to unpredictable thread ordering
 - Usually paired with some kind of state variable
 - A counter in the case of a semaphore
 - State variable is protected by its own spinlock
 - Must ensure the state variable is modified together with the wait channel
 - Thread must therefore always be holding at least one of the two locks
 - Must also ensure that all locks are released before going to sleep
 - Locks must be held to wake up a sleeping thread
 - Therefore, if *wchan_sleep* acquires *wc_lock* on its own, then it is difficult to:
 1. Release the state variable's lock
 2. while holding *wc_lock* and
 3. before the thread goes to sleep

- **Semaphore Review**

- A semaphore counter should never be negative (not even temporarily)
- P() must therefore block when the counter value is 0
- Important details of **our** P() implementation:

```
1  P(struct semaphore* sem) {
2      spinlock_acquire(&sem->sem_lock);
3      while (sem->sem_count == 0) {
4          wchan_lock(sem->sem_wchan);
5          spinlock_release(&sem->sem_lock);
6          wchan_sleep(sem->sem_wchan);
7          spinlock_acquire(&sem->sem_lock);
8      }
9      sem->sem_count--;
10     spinlock_release(&sem->sem_lock);
11 }
```

- Does not provide FIFO ordering. Example:
 - Semaphore count is 1 after waking up a sleeping thread
 - Another thread can decrement it back to 0 before the previously sleeping thread gets to run
- Must check the semaphore count after waking up
 - *While* loop instead of an *if* statement at line 3
- *sem_lock* must be released after acquiring the wait channel lock and before *wchan_sleep*